Weekly Readings Archive for GammaTauAI

Week	Paper
Week 1	Multi-Frame Self-Supervised Depth with Transformers
Week 2	Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor
Week 3	Collective Intelligence for Deep Learning: A Survey of Recent Developments
Week 3	One Policy to Control Them All: Shared Modular Policies for Agent-Agnostic Control
Week 4	Top-Down Synthesis for Library Learning
Week 4	DreamCoder: Growing generalizable, interpretable knowledge with wake-sleep Bayesian program learning
Week 5	Automated Antenna Design via Domain Knowledge-Informed Reinforcement Learning and Imitation Learning
Week 6	Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm
Week 6	Transformer Feed-Forward Layers Are Key-Value Memories
Week 7	The Suprising Creativity of Digital Evolution: A Collection of Anecdotes from the Evolutionary Computation and Artificial Life Research Communities
Week 8	The Alignment Problem from a Deep Learning Perspective
Week 8	GFNs
Week 9	Presentation from Federico Cassano
Week 10	Presentation from Jacob (Jake) Ginesin
Week 11	A Path Towards Autonomous Machine Intelligence
Week 12	Presentation from Caden Juang
Week 13	Metallaxis-FL: Mutation-based Fault Localization
Week 14	Discussion of spoken language and affects on cognition
Week 15	All You Ever Wanted to Know About Dynamic Taint Analysis and Forward Symbolic Execution (but might have been afraid to ask)
Week 15	Static versus dynamic analysis—an illusory distinction?)