Weekly Readings Archive for GammaTauAI

Week	Paper
Week 1	Multi-Frame Self-Supervised Depth with Transformers
Week 2	Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor
Week 3	Collective Intelligence for Deep Learning: A Survey of Recent Developments
Week 3	One Policy to Control Them All: Shared Modular Policies for Agent-Agnostic Control
Week 4	Top-Down Synthesis for Library Learning
Week 4	DreamCoder: Growing generalizable, interpretable knowledge with wake-sleep Bayesian program learning
Week 5	Automated Antenna Design via Domain Knowledge-Informed Reinforcement Learning and Imitation Learning
Week 6	Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm
Week 6	Transformer Feed-Forward Layers Are Key-Value Memories
Week 7	The Suprising Creativity of Digital Evolution: A Collection of Anecdotes from the Evolutionary Computation and Artificial Life Research Communities
Week 8	The Alignment Problem from a Deep Learning Perspective
Week 8	GFNs
Week 9	Presentation from Federico Cassano
Week 10	Presentation from Jacob (Jake) Ginesin
Week 11	A Path Towards Autonomous Machine Intelligence
Week 12	Presentation from Caden Juang
Week 13	Metallaxis-FL: Mutation-based Fault Localization
Week 14	Discussion of spoken language and affects on cognition
Week 15	All You Ever Wanted to Know About Dynamic Taint Analysis and Forward Symbolic Execution (but might have been afraid to ask)
Week 15	Static versus dynamic analysis—an illusory distinction?

Week	Paper
Week 16	PROMPTBREEDER: SELF-REFERENTIAL SELF-IMPROVEMENT VIA PROMPT EVOLUTION
Week 17	Presentation from Gio on Differential Dynamic Logic for Hybrid Systems
Week 18	How to Evaluate Blame for Gradual Types
Week 19	Of Non-Linearity and Commutativity in BERT
Week 20	Evolution through Large Models
Week 21	MACHINE LEARNING AND INFORMATION THEORY CONCEPTS TOWARDS AN AI MATHEMATICIAN
Week 22	Training Verifiers to Solve Math Word Problems
Week 22	Let's Verify Step by Step
Week 23	ZeRO: Memory Optimizations Toward Training Trillion Parameter Models
Week 24	KAN: Kolmogorov-Arnold Networks
Week 25	Q-Probe: A Lightweight Approach to Reward Maximization for Language Models
Week 26	Activation Steering for Robust Type Prediction in CodeLLMs. Guest talk from Francessa Lucchetti
Week 27	egg: Fast and Extensible Equality Saturation
Week 28	Transformer Layers as Painters
Week 29	Quiet-STaR: Language Models Can Teach Themselves to Think Before Speaking